

Convolutional Neural Networks

Yao Zhang

Mostly based on Thomas Hofmann's lecture in ETH

<https://zhims.github.io/datascience.html>

Nov 12, 2019

Integral Operators

Definition 1 (Integral operator)

A transform T expressible with the kernel H and $t_1, t_2 \in \mathbb{R} \cup \{-\infty, \infty\}$ such that for any function f (for which Tf exists)

$$(Tf)(u) = \int_{t_1}^{t_2} H(u, t) f(t) dt \quad (1)$$

is called an integral operator.

Example 1 (Fourier transform)

$$(\mathcal{F}f)(u) \triangleq \int_{-\infty}^{\infty} e^{-2\pi i t u} f(t) dt \quad (2)$$

Definition 2 (Convolution)

Given two functions f, h , their convolution is defined as

$$(f * h)(u) \triangleq \int_{-\infty}^{\infty} h(u - t) f(t) dt = \int_{-\infty}^{\infty} f(u - t) h(t) dt \quad (3)$$

Remark 1

- 1 *integral operator with kernel $H(u, t) = h(u - t)$*
- 2 *shift-invariant* as $H(u - s, t - s) = h(u - t) = H(u, t) \quad (\forall s)$
- 3 *convolution operator is commutative*
- 4 *existence depends on properties of f, h*
- 5 *typical use $f = \text{signal}$, $h = \text{fast decaying kernel function}$*

Linear Time-Invariant Transforms

Definition 3 (Linear transform)

T is linear, if for all functions f, g and the scalars α, β ,

$$T(\alpha f + \beta g) = \alpha Tf + \beta Tg \quad (4)$$

Definition 4 (Translation invariant transform)

T is translation (or shift) invariant, if for any f and scalar τ ,

$$f_\tau(t) \triangleq f(t + \tau), \quad (Tf_\tau)(t) \triangleq (Tf)(t + \tau) \quad (5)$$

Theorem 1

*Any linear, translation-invariant transformation T can be written as **convolution** with a suitable h .*

Signal Processing with Neural Networks

- ① Transforms in deep networks: **linear** + simple non-linearity
- ② Many signals (audio, image, etc.) obey **translation invariance** \Rightarrow **invariant** feature maps: shift in input = shift in feature map

1 + 2 in above:

- ① \Rightarrow learn convolutions, not (full connectivity) weight matrices
- ② \Rightarrow **convolutional layers** for signal processing

Discrete Convolutions

For all practical purposes: signal are sampled, i.e. discrete.

Definition 5 (Discrete convolution (1-D))

For $f, h : \mathbb{Z} \rightarrow \mathbb{R}$, we can define the discrete convolution via

$$(f * h)[u] \triangleq \sum_{t=-\infty}^{\infty} f[t] h[u - t] \quad (6)$$

Remark 2

① *use of rectangular brackets to suggest "arrays"*

② *2D case:*

content... (7)

③ *typical: h with finite support (window size)*

Discrete Convolutions: Example

Example 2

Small Gaussian kernel with support $[-2 : 2] \subset \mathbb{Z}$

$$h[t] = \frac{1}{16} \begin{cases} 6 & t = 0 \\ 4 & |t| = 1 \\ 1 & |t| = 2 \\ 0 & \text{otherwise} \end{cases} \quad (8)$$

Consequence: convolution sum can be truncated:

$$\begin{aligned} (f * h)[u] &= \sum_{t=u-2}^{u+2} f[t] h[u-t] = \sum_{t=-2}^2 h[t] f[u-t] \\ &= \frac{6f[u] + 4f[u-1] + 4f[u+1] + f[u-2] + f[u+2]}{16} \end{aligned} \quad (9)$$

Definition 6 (Discrete cross-correlation)

Let $f, h : \mathbb{Z} \rightarrow \mathbb{R}$, then

$$(h * f)[u] \triangleq \sum_{t=-\infty}^{\infty} h[t] f[u + t] \quad (10)$$

Remark 3

- ① Def. 6 also called a "sliding inner product", $u + t$ instead of $u - t$
- ② note that cross-correlation and convolution are closely related:

$$\begin{aligned}(h * f)[u] &= \sum_{t=-\infty}^{\infty} h[t] f[u + t] \\&= (h * f)[u] = \sum_{t=-\infty}^{\infty} h[-t] f[u - t] \\&= (\bar{h} * f)[u] \\&= (f * \bar{h})[u]\end{aligned}\tag{11}$$

where $\bar{h}[t] \triangleq h[-t]$.

Only difference: kernel flipped over, but not non-commutative.

Convolution via Matrices

- ① In practice: signal f and kernel h have finite support
- ② Without loss of generality (w.l.o.g) $f[t] = 0$ for $t \notin [1 : n]$, $h[t] = 0$ for $t \notin [1 : m]$
- ③ We can think of f and h as vectors and define:

$$(f * h) = \underbrace{\begin{pmatrix} h_1 & 0 & 0 & 0 & \cdots & 0 & 0 \\ h_2 & h_1 & 0 & 0 & \cdots & 0 & 0 \\ h_3 & h_2 & h_1 & 0 & \cdots & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & \cdots & h_m & h_{m-1} \\ 0 & 0 & 0 & 0 & \cdots & 0 & h_m \end{pmatrix}}_{\triangleq H_n^h \in \mathbb{R}^{(n+m-1) \times n}} \begin{pmatrix} f_1 \\ f_2 \\ f_3 \\ \vdots \\ f_{n-1} \\ f_n \end{pmatrix} \quad (12)$$

Toeplitz Matrix

Definition 7 (Toeplitz matrix)

A matrix $H \in \mathbb{R}^{k \times n}$ is a Toeplitz matrix, if there exists $n + k - 1$ numbers c_l ($l \in [-(n-1) : (k-1)] \subset \mathbb{Z}$) such that

$$H_{ij} = c_{i-j} \quad (13)$$

Remark 4

- ① *in plain English, all **NW-SE** diagonals are constant*
- ② *if $m \ll n$: additional sparseness (band matrix of width m)*
- ③ *H_n^h has only m degrees of freedom*
- ④ *locality (sparseness $m \ll n$) and weight sharing (kernel)*

Sparse Connectivity




Figure 1: Sparse vs dense connectivity

Convolutions in Higher Dimensions

Generalize concept of convolution to:

- 1 2D: e.g. images, spectrograms
- 2 3D: e.g. color or multi-spectral images, voxel images, video
- 3 or even higher dimensions

Replace vector by:

- 1 matrices or fields (e.g. in discrete case)

$$(F * G)[i, j] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} F[i - k, j - l] \cdot G[k, l] \quad (14)$$

- 2 tensors: for 3D and higher

Convolutional Layers: Border Handling

Different options for border handling:

- ① our definition: padding with zeros = **same padding**
- ② only retain values from windows fully contained in support of signal f = **valid padding**




Figure 2:

Convolutional Layers: Layout

- ① Convolved signal **inherits** topology of original signal
- ② Hence: units in a convolutional layer are typically arranged on the same grid (1D, 2D, 3D,...)

Convolutional Layers: Backpropagation

Exploit **structural sparseness** in computing $\frac{\partial x_i^l}{\partial x_j^{l-1}}$:

- ① receptive field of $x_i^l : \mathcal{I}_i^l \triangleq \{j : W_{ij}^l \neq 0\}$, where W^l is the Toeplitz matrix of the convolution
- ② obviously $\frac{\partial x_i^l}{\partial x_j^{l-1}} = 0$ for $j \notin \mathcal{I}_i^l$

Weight sharing in computing $\frac{\partial \mathcal{R}}{\partial h_j^l}$, where h_j^l is a kernel weight

$$\frac{\partial \mathcal{R}}{\partial h_j^l} = \sum_i \frac{\partial \mathcal{R}}{\partial x_i^l} \frac{\partial x_i^l}{\partial h_j^l} \quad (15)$$

Weight is re-used for every unit within target layer \Rightarrow additive combination of derivatives in chain rule.

Efficient Computations of Convolutional Activities

FFT (**Fast Fourier Transform**): compute convolutions fast(er).

- 1 Fourier transform of signal $f \rightarrow (\mathcal{F}f)$ and kernel $h \rightarrow (\mathcal{F}h)$
- 2 pointwise multiplication and inverse Fourier transform:

$$(f * h) = \mathcal{F}^{-1}((\mathcal{F}f) \cdot (\mathcal{F}h)) \quad (16)$$

- 3 FFT: signal of length n , can be done in $O(n \log n)$
- 4 pays off, if many channels (amortizes computation of $\mathcal{F}f$)
- 5 small kernels ($m < \log n$): favor time / space domain

Convolutional Layers: Stages

- ① **Non-linearities**: detector stage. As always: scalar non-linearities (activation function)
- ② **Pooling** stage: locally combine activities

Most frequently used pooling function: **max pooling**.

Definition 8 (Max Pooling)

Define window size r (e.g. 3 or 3×3), then

$$\begin{aligned} 1D : \quad x_i^{\max} &= \max \{x_{i+k} : 0 \leq k < r\}, \\ 2D : \quad x_{ij}^{\max} &= \max \{x_{i+k,j+l} : 0 \leq k, l < r\} \end{aligned} \tag{17}$$

Remark 5

- ① *maximum over a small patch of units*
- ② *other functions are possible: average, soft-maximization*

Max-pooling

Max-pooling: invariance

- ① set of invertible transformations \mathcal{T} : group w.r.t composition
- ② \mathcal{T} -invariance through maximization $f_{\mathcal{T}}(x) \triangleq \max_{\tau \in \mathcal{T}} f(\tau x)$

Proposition 1

$f_{\mathcal{T}}$ is invariant under $\tau \in \mathcal{T}$.

Proof.

$$f_{\mathcal{T}}(\tau x) = \max_{\rho \in \mathcal{T}} f(\rho(\tau x)) = \max_{\rho \in \mathcal{T}} (f(\rho \circ \tau)x) = \max_{\sigma \in \mathcal{T}} f(\sigma x) \quad (18)$$

as $\forall \sigma, \sigma = \rho \circ \tau$ with $\rho = \sigma \circ \tau^{-1}$. □

Sub-Sampling(also known as (aka) Strides)

- ① often, it is desirable to reduce the size of feature maps
- ② **sub-sampling**: reduce temporal/spatial resolution. Often: combined with (max-)pooling (aka. stride)
- ③ example: max-pool, filter 2×2 , stride 2×2
- ④ disadvantage: loss of information

Learn multiple convolution kernel (or filters) = multiple **channels**:

- ① typically: all channels use same window size
- ② channels form additional dimension for next layer (e.g. 2D signal \times channels = 3D tensor)
- ③ number of channels: design parameter

Convolutional Layers: Animation

<http://cs231n.github.io/assets/conv-demo/index.html>

Convolutional Layers for Vision

Note that kernels (across channels) form a linear map:

$$h : \mathbb{R}^{r^2 \times d} \rightarrow \mathbb{R}^k \quad (19)$$

where $r \times r$ is the window size and d is the depth.




Figure 3: convolutional layers for vision

Convolutional Networks: ConvNets

Convolutional networks: **multiple, stacked feature maps**

$$\underbrace{y[r]}_{r\text{-th channel}}[s, t] = \sum_u \sum_{\Delta s, \Delta t} \underbrace{w[r, u][\Delta s, \Delta t]}_{\text{parameters}} \underbrace{x[u]}_{u\text{-th channel}}[s + \Delta s, t + \Delta t] \quad (20)$$

- ❶ x, y tensor, 3-rd order
- ❷ number of parameters:

$$\underbrace{\#r \cdot \#u}_{\text{fully connected}} \cdot \underbrace{\#\Delta s \cdot \#\Delta t}_{\text{window size}} \quad (21)$$

- ❸ pointwise non-linearities (e.g. ReLU)
- ❹ interleaved with: pooling (e.g. max, average)
- ❺ optionally: downsampling (use of strides)

Convolutional Pyramid

Typical use of convolution in vision: sequence of convolutions that

- 1 **reduce** spatial dimensions (sub-sampling)
- 2 **increase** number of channels

⇒ smaller, but more feature maps.




Figure 4: Architecture of LeNet-5, a convolutional neural network, here for digits recognition. Each plan is a feature map, i.e. a set of units whose weights are constrained to be identical.

- ① C1/S2: 6 channels, 5×5 kernels, 2×2 sub (4704 units)
- ② C3/S4: 16 channels, 6×6 kernels, 2×2 sub (1600 units)
- ③ C5: 120 channels, F6: fully-connected
- ④ output: Gaussian noise model (squared loss)




Figure 5: AlexNet architecture

- ① **Pyramidal architecture:** reduce spatial resolution, increase channels with depth
- ② Challenge: many channels (width) + large windows + depth
- ③ Number of parameters
 - ① 384 to 384 channels with 3×3 windows: > 1.3 M
 - ② $13 \times 13 \times 384$ tensor to 4096, fully connected: > 265 M

Deep ConvNets: Key Challenges

- ① avoid blow-up of model size (e.g. # parameters)
- ② preserve computational efficiency of learning (e.g. gradients)
- ③ allow for large depth (as it is known to be a plus)
- ④ allow for sufficient width (as it is known to be a plus, too)

Very Deep Convolutional Networks: VGG




Figure 6: VGG 16

- 1 use very small receptive fields (maximally 3×3)
- 2 avoid downsampling/pooling
- 3 stacking small receptive fields: more depth, fewer parameters
- 4 example: $3 \cdot (3 \times 3) = 27 < 49(7 \times 7)$

Inception Module: 1×1 Convolution

Many channels needed for high accuracy, typically $k \sim 200 - 1000$ (e.g. AlexNet: 2×192).


Observation (motivated by Arora et al, 2013): when convolving, dimension reduction across channels may be acceptable.

Dimension reduction: m channels of a $1 \times 1 \times k$ convolution $m \leq k$:


$$x_{ij}^+ = \sigma(Wx_{ij}), \quad W \in \mathbb{R}^{m \times k} \quad (22)$$

- ① 1×1 convolution = no convolution
- ② inception module (Szegedy et al.)
- ③ network within a network (Lin et al.)
- ④ i.e. W is shared for all (i, j) (translation invariance)

Inception Module: Mixing



(a) naive version



(b)

Figure 7: Inception Module

Instead of fixed window size convolution: **mix** 1×1 with 3×3 and 5×5 , max-polling. Use 1×1 convolutions for dimension reduction before convolving with large kernels.

Google Inception Network

Very deep network: many inception modules (green boxes: concatenation points). Additional trick: connect softmax layer (and loss) at intermediate stages (yellow boxes) \Rightarrow gradient shortcuts.

Google Inception Network

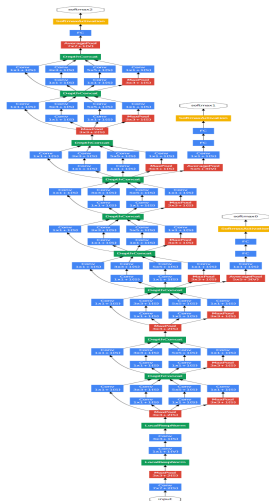


Figure 8: Google inception networks

Residual Networks: ResNets




Figure 9: Residual Networks module

- 1 learn changes to the **identity map** (aka. shortcut connections)
- 2 use small filters (VGG), use dimension reduction (inception)
- 3 reach depth of 100 + layers (+ increase accuracy + trainable)

Thank you all of you! –Yao